

CANNs: Continuous Attractor Neural Networks Toolkit with ASA for Attractor Structure Analysis

Sichao He¹, Aiersi Tuerhong², Shangjun She³, Tianhao Chu³, Yuling Wu¹, Junfeng Zuo¹, and Si Wu^{1, 3, 4, 5, 6}

1 Peking-Tsinghua Center for Life Sciences, Academy for Advanced Interdisciplinary Studies, Peking University, Beijing, China **2** College of Mathematics and Statistics, Chongqing University, Chongqing, China **3** School of Psychological and Cognitive Sciences, Peking University, Beijing, China **4** School of Psychology and Cognitive Sciences, Peking University, Beijing, China **5** PKU-IDG/McGovern Institute for Brain Research, Peking University, Beijing, China **6** Center of Quantitative Biology, Peking University, Beijing, China

DOI:

Software

- [Review ↗](#)
- [Repository ↗](#)
- [Archive ↗](#)

Submitted:

Published:

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC-BY](#)).

Summary

CANNs (Continuous Attractor Neural Networks toolkit) is a Python library built on BrainPy, a powerful framework for brain dynamics programming. It streamlines experimentation with continuous attractor neural networks and related brain-inspired models. The library delivers ready-to-use models, task generators, analysis tools, and pipelines—enabling neuroscience and AI researchers to move quickly from ideas to reproducible simulations.

Statement of need

Continuous Attractor Neural Networks (CANNs) provide a theoretical framework for understanding how the brain encodes continuous variables—such as spatial position, head direction, and movement trajectories—through stable neural activity patterns. These models explain key phenomena in hippocampal place cells (O’Keefe and Dostrovsky 1971), entorhinal grid cells (Hafting et al. 2005), and head direction systems (Taube, Müller, and Ranck 1990). Despite their importance, CANN research suffers from fragmentation: researchers implement models from scratch, use incompatible codebases, and face significant reproducibility barriers. This lack of standardization slows progress and creates steep learning curves for newcomers.

CANNs addresses this gap by providing a unified Python toolkit built on BrainPy (Wang et al. 2023). It delivers: (1) standardized implementations of CANN and related brain-inspired models, including canonical frameworks (Wu-Amari-Wong (Amari 1977; Wu, Hamaguchi, and Amari 2008), adaptation-augmented CANNs (Mi et al. 2014; Li, Chu, and Wu 2025), grid cell networks (Burak and Fiete 2009)), alongside additional attractor architectures; (2) integrated task generation, simulation, and analysis pipelines; and (3) high-performance computation via JAX JIT compilation and optional Rust acceleration. By standardizing workflows—analogous to Hugging Face Transformers in deep learning—this library accelerates reproducible research and lowers barriers for computational neuroscientists, AI engineers, and students exploring attractor dynamics.

Software design

Layer hierarchy of the CANNs library showing five levels: Application (Pipeline orchestration), Functional (Task, Trainer, Analyzer, Utils modules), Core Models (CANN implementations), Foundation (BrainPy/JAX and Rust FFI backends), and Hardware (CPU/GPU/TPU support).

The CANNs library follows a modular architecture (section) guided by two core principles: **separation of concerns** and **extensibility through base classes**. The design separates functional responsibilities into five independent modules: (1) **Models** (`canns.models`) define neural network dynamics; (2) **Tasks** (`canns.task`) generate experimental paradigms and input data; (3) **Analyzers** (`canns.analyzer`) provide visualization and analysis tools; (4) **Trainers** (`canns.trainer`) implement learning rules for brain-inspired models; and (5) **Pipeline** (`canns.pipeline`) orchestrates complete experimental workflows.

Each module focuses on a single responsibility—models don’t generate input data, tasks don’t analyze results, and analyzers don’t modify parameters. This separation ensures maintainability, testability, and extensibility. All major components inherit from abstract base classes (`BasicModel`, `BrainInspiredModel`, `Trainer`) that define standard interfaces, enabling users to create custom implementations that seamlessly integrate with the built-in ecosystem.

The library supports four distinct research workflows: (1) CANN modeling and simulation for studying attractor dynamics; (2) data analysis for processing experimental neural recordings; (3) brain-inspired learning with biologically plausible plasticity rules (Hebbian, STDP, BCM); and (4) end-to-end pipelines for automated parameter sweeps and reproducible experiments. All models inherit from BrainPy’s `DynamicalSystem` base class (Wang et al. 2023), leveraging JAX’s JIT compilation for GPU/TPU acceleration while maintaining simple Python APIs. For operations where Python overhead is significant, the companion `canns-lib` Rust library provides optional accelerated backends—notably achieving 400× speedup for spatial navigation tasks and 1.13-1.82× speedup for topological data analysis—without requiring code structure changes.

Related Works

Computational neuroscience has established simulators for spiking neural networks, including NEST (Gewaltig and Diesmann 2007) for large-scale network models, Brian 2 (Stimberg, Brette, and Goodman 2019) for rapid prototyping with user-friendly syntax, and NEURON (Hines and Carnevale 1997) for detailed biophysical simulations. However, these general-purpose tools lack specialized support for continuous attractor dynamics and require significant implementation effort for CANN-specific workflows. While individual CANN implementations exist (e.g., `cann_base` on GitHub), they remain fragmented, lab-specific codebases without standardized APIs, comprehensive task generation, or analysis tools—paralleling the pre-standardization era of deep learning before frameworks like Hugging Face Transformers (Wolf et al. 2020) unified model definitions and usage patterns.

CANNs builds upon BrainPy (Wang et al. 2023), a modern brain dynamics framework leveraging JAX (Bradbury et al. 2018) for JIT compilation and autodifferentiation. BrainPy provides the foundational infrastructure—state management, time stepping, and GPU/TPU acceleration—that CANNs extends with CANN-specific abstractions: standardized model implementations, task-generation APIs, analysis pipelines, and Rust-accelerated backends. This layered approach mirrors successful deep learning ecosystems: PyTorch (Ansel et al. 2024) and TensorFlow (Abadi et al. 2015) provide low-level

tensor operations, while domain-specific libraries (e.g., Transformers [(Wolf et al. 2020)] for NLP, torchvision [torchvision2016] for computer vision) deliver standardized high-level components. CANNs fills this role for continuous attractor research, reducing fragmentation and accelerating reproducible science.

AI usage disclosure

AI-assisted tools were used during the development of this software in limited, auxiliary capacities. Sourcery AI was employed for automated code quality reviews and refactoring suggestions in pull requests. Large language models (including Claude and ChatGPT) assisted with documentation writing, including docstring generation, tutorial content drafting, and technical writing refinement. However, all core library code—including model implementations, task generators, analyzers, and algorithmic components—was written by human developers. All AI-generated content was reviewed, validated, and edited by the authors to ensure technical accuracy and consistency with project standards. The software architecture, design decisions, and scientific contributions represent original human intellectual work.

Acknowledgements

We thank Aiersi Tuerhong and Shangjun She for their collaborative development contributions to the library implementation. We are grateful to Tianhao Chu, Yuling Wu, and Junfeng Zuo for valuable discussions, feedback, and guidance throughout the project development. We especially thank Si Wu for overall scientific guidance and mentorship. We acknowledge the BrainPy development team for providing the foundational framework upon which this library is built, and the broader open-source community for tools and libraries that enabled this work.

References

- Abadi, Martín, Ashish Agarwal, Paul Barham, Eugene Brevdo, Zhifeng Chen, Craig Citro, Greg S. Corrado, et al. 2015. *TensorFlow, Large-scale machine learning on heterogeneous systems*. <https://doi.org/10.5281/zenodo.4724125>.
- Amari, Shun-ichi. 1977. “Dynamics of Pattern Formation in Lateral-Inhibition Type Neural Fields.” *Biological Cybernetics* 27 (2). Springer:77–87. <https://doi.org/10.1007/BF00337259>.
- Ansel, Jason, Edward Yang, Horace He, Natalia Gimelshein, Animesh Jain, Michael Voznesensky, Bin Bao, et al. 2024. “PyTorch 2: Faster Machine Learning Through Dynamic Python Bytecode Transformation and Graph Compilation.” In *29th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 2 (Asplos ’24)*. ACM. <https://doi.org/10.1145/3620665.3640366>.
- Bradbury, James, Roy Frostig, Peter Hawkins, Matthew James Johnson, Chris Leary, Dougal Maclaurin, George Necula, et al. 2018. *JAX: Composable Transformations of Python+NumPy Programs* (version 0.3.13). <http://github.com/jax-ml/jax>.
- Burak, Yoram, and Ila R Fiete. 2009. “Accurate Path Integration in Continuous Attractor Network Models of Grid Cells.” *PLoS Computational Biology* 5 (2). Public Library of Science San Francisco, USA:e1000291.

Gewaltig, Marc-Oliver, and Markus Diesmann. 2007. “NEST (Neural Simulation Tool).” *Scholarpedia* 2 (4):1430.

Hafting, Torkel, Marianne Fyhn, Sturla Molden, May-Britt Moser, and Edvard I Moser. 2005. “Microstructure of a Spatial Map in the Entorhinal Cortex.” *Nature* 436 (7052). Nature Publishing Group UK London:801–6. <https://doi.org/10.1038/nature03721>.

Hines, Michael L, and Nicholas T Carnevale. 1997. “The Neuron Simulation Environment.” *Neural Computation* 9 (6). MIT Press:1179–1209.

Li, Yujun, Tianhao Chu, and Si Wu. 2025. “Dynamics of Continuous Attractor Neural Networks with Spike Frequency Adaptation.” *Neural Computation* 37 (6). MIT Press 255 Main Street, 9th Floor, Cambridge, Massachusetts 02142, USA ...:1057–1101. https://doi.org/10.1162/neco_a_01757.

Mi, Yuanyuan, CC Fung, KY Wong, and Si Wu. 2014. “Spike Frequency Adaptation Implements Anticipative Tracking in Continuous Attractor Neural Networks.” *Advances in Neural Information Processing Systems* 27.

O’Keefe, John, and Jonathan Dostrovsky. 1971. “The Hippocampus as a Spatial Map: Preliminary Evidence from Unit Activity in the Freely-Moving Rat.” *Brain Research*. Elsevier Science. [https://doi.org/10.1016/0006-8993\(71\)90358-1](https://doi.org/10.1016/0006-8993(71)90358-1).

Stimberg, Marcel, Romain Brette, and Dan FM Goodman. 2019. “Brian 2, an Intuitive and Efficient Neural Simulator.” *eLife* 8. eLife Sciences Publications, Ltd:e47314.

Taube, Jeffrey S, Robert U Muller, and James B Ranck. 1990. “Head-Direction Cells Recorded from the Postsubiculum in Freely Moving Rats. I. Description and Quantitative Analysis.” *Journal of Neuroscience* 10 (2). Society for Neuroscience:420–35. <https://doi.org/10.1523/JNEUROSCI.10-02-00420.1990>.

Wang, Chaoming, Tianqiu Zhang, Xiaoyu Chen, Sichao He, Shangyang Li, and Si Wu. 2023. “BrainPy, a Flexible, Integrative, Efficient, and Extensible Framework for General-Purpose Brain Dynamics Programming.” Edited by Marcel Stimberg. *eLife* 12 (December). eLife Sciences Publications, Ltd:e86365. <https://doi.org/10.7554/eLife.86365>.

Wolf, Thomas, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierrick Cistac, et al. 2020. “Transformers: State-of-the-Art Natural Language Processing.” In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, 38–45. Online: Association for Computational Linguistics. <https://www.aclweb.org/anthology/2020.emnlp-demos.6>.

Wu, Si, Kosuke Hamaguchi, and Shun-ichi Amari. 2008. “Dynamics and Computation of Continuous Attractors.” *Neural Computation* 20 (4). MIT Press One Rogers Street, Cambridge, MA 02142-1209, USA journals-info ...:994–1025. <https://doi.org/10.1162/neco.2008.10-06-378>.